

# Simplified and Generalized Masked Diffusion for Discrete Data

Jixin Shi\*, Kehang Han\*, Zhe Wang, Arnaud Doucet, Michalis K. Titsias  
Google DeepMind

## Abstract

They call their model MD4:  
Masked Discrete Diffusion  
for Discrete Data

Masked diffusion models have shown great promise in generating discrete data. However, the autoregressive nature of masked diffusion models makes them difficult to train. In this paper, we propose a simplified and generalized framework that unlocks the full potential of masked diffusion models. We show that the continuous-time variational objective of masked diffusion models is a simple weighted integral of cross-entropy losses. Our framework also enables training generalized masked diffusion models with state-dependent masking schedules. When evaluated by perplexity, our models trained on OpenWebText surpass prior diffusion language models at GPT-2 scale and demonstrate superior performance on 4 out of 5 zero-shot language modeling tasks. Furthermore, our models vastly outperform previous discrete diffusion models on pixel-level image modeling, achieving 2.75 (CIFAR-10) and 3.40 (ImageNet 64×64) bits per dimension that are better than autoregressive models of similar sizes. Our code is available at <https://github.com/google-deepmind/md4>.

# Motivation: Why do we need Discrete Diffusion?

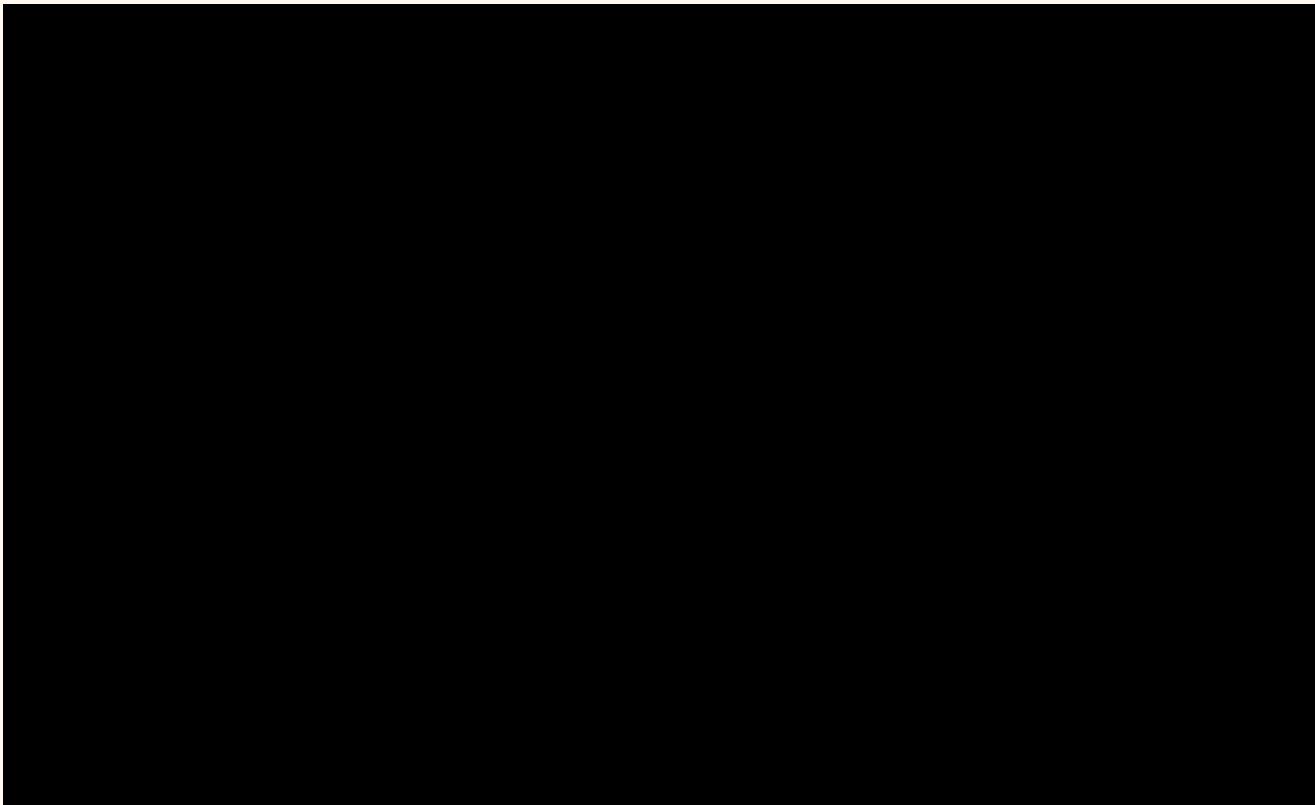
AI currently handles two fundamentally different data types

Standard (Gaussian) Diffusion	The Discrete Challenge
<b>Used for:</b> Continuous data (Images, Audio)	<b>Targeting:</b> Discrete data (Text, Code, DNA sequences)
<b>Mechanism:</b> Add Gaussian noise.	<b>Constraint:</b> Data is categorical.
<b>Why it works there:</b> An image pixel is continuous data. A "slightly noisy pixel" is still a valid concept.	<b>The Problem:</b> There is no such thing as "half a word" or "Word A + 0.1 * Word B".
<b>The Failure Mode on Text:</b> Adding Gaussian noise to the token embedding for "Cat" creates a nonsense vector that does not map to <i>any</i> word in the vocabulary dictionary.	<b>The Goal:</b> We want the benefits of diffusion (iterative refinement, holistic/non-causal generation) but mathematically adapted specifically for a discrete vocabulary.

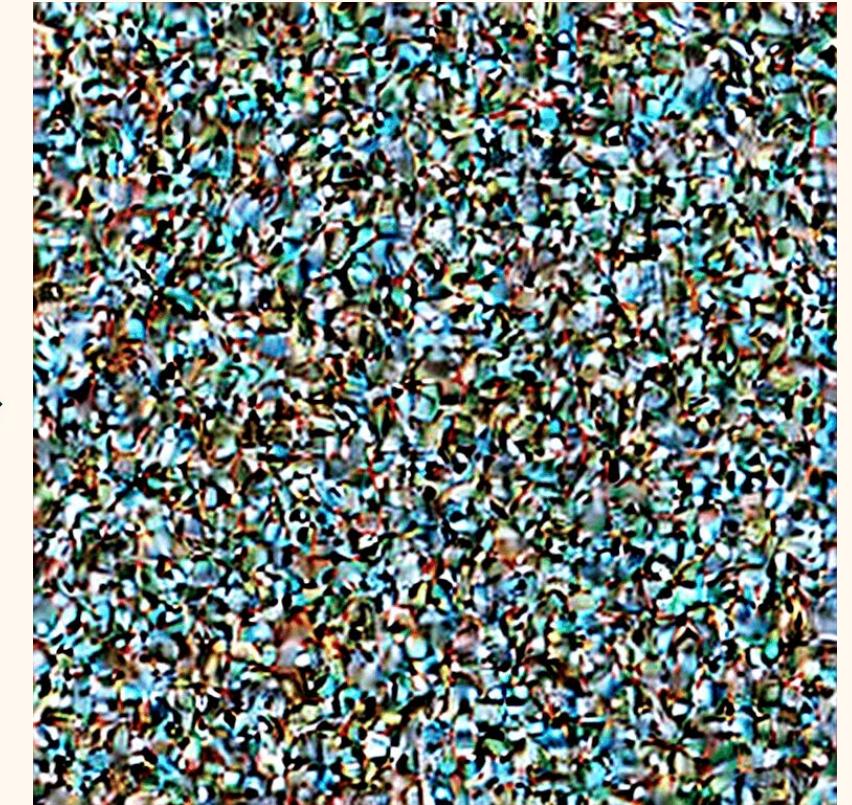
Solution: Instead of noise, we add masks!

# What does discrete diffusion look like?

Starting from a fully masked state (Discrete)

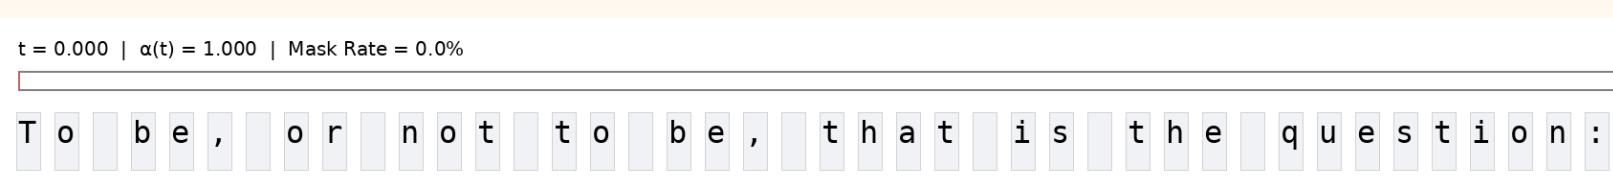


Starting from noise (continuous)



# The simplified masked Forward Process (1/2)

- **$t = 0$  (No noise),  $t = 1$  (All masks),  $x_0$  (original data)**
- **The General Framework (D3PM<sup>\*</sup>):** Transition Matrix  $Q_t$  representing the  $P(i \rightarrow j)$  of a token transitioning from state  $i$  to state  $j$ .
- **The Problem:** Calculating a full  $Q_t$  needs  $K^2$  entries. For large vocabulary sizes (e.g.  $K = 50k$ ) this becomes prohibitive.
- **The Simplification - Absorbing State.**
- **The Rule:** A token has only two possible states:
  - **Stay  $x_0$ :** It remains the original token.
  - **Become [MASK]:** It is corrupted into a special mask token.
- **Absorbing Property:** Once a token becomes [MASK], it stays [MASK] for the rest of the forward process.



# The simplified masked Forward Process (2/2)

- $t = 0$  (No noise),  $t = 1$  (All masks),  $x_0$  (original data)

**The Marginal Distribution** simplifies into a direct formula for the state at time t:

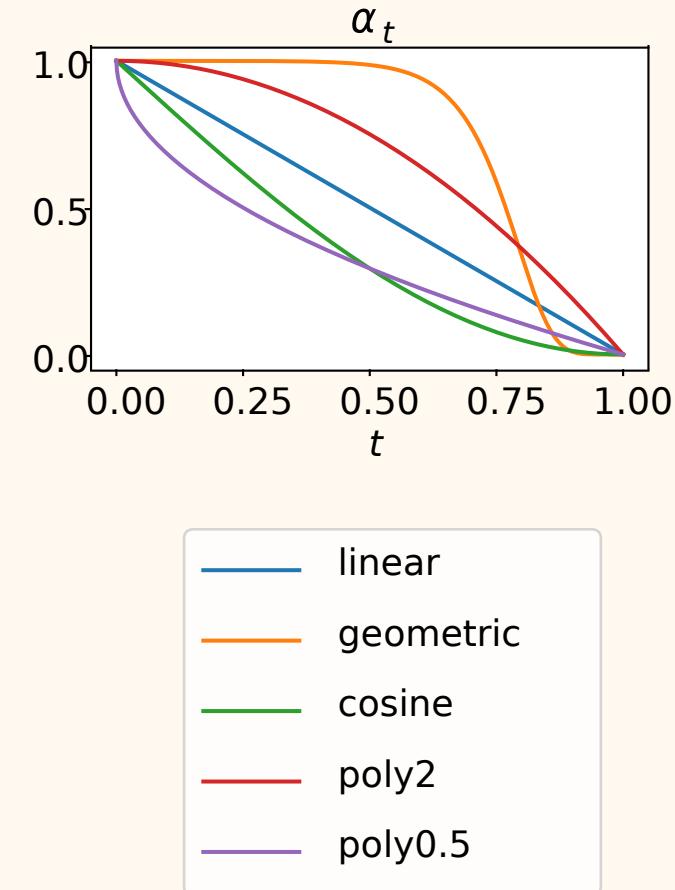
$$q(x_t | x_0) = \alpha_t \cdot \delta_{x_0} + (1 - \alpha_t) \cdot \delta_{\text{Mask}}$$

**Interpretation:**

- At any time t, the token is either the original  $x_0$  (with prob  $\alpha_t$ ) or the [MASK] (with prob  $1 - \alpha_t$ ).
- $\alpha_t$  is a schedule that we can choose.  $t = 1 \rightarrow \alpha_t = 0$

**The final Forward Process:** This easily gives us now the formula for any two times  $0 \leq s \leq t \leq 1$ :

$$q(x_t | x_s) = \frac{\alpha_t}{\alpha_s} \cdot \delta_{x_s} + \left(1 - \frac{\alpha_t}{\alpha_s}\right) \cdot \delta_{\text{Mask}}$$



# The Reverse Process (Posterior)

- $t = 0$  (No noise),  $t = 1$  (All masks),  $x_0$  (original data),  $0 \leq s \leq t \leq 1$

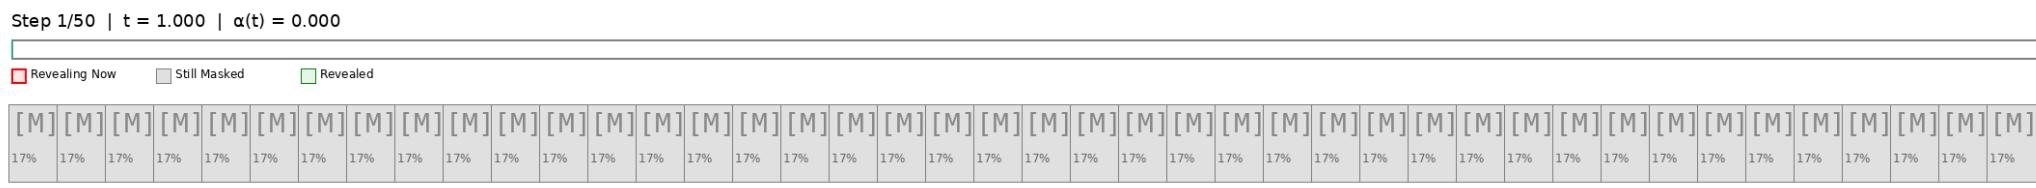
We look at the posterior  $q(x_s|x_t, x_0)$ .

**“Absorbing” in the forward process -> deterministic in the reverse process:**

1. If  $x_t$  is unmasked, then  $x_s = x_t$ . (Masks cannot appear out of nowhere in the backward.)
  2. If  $x_t$  is [MASK]: We have a chance to reveal the token  $x_0$ .

Probability for a token to get revealed in this specific step  $t \rightarrow s$ :

$$P(\text{reveal}) = \frac{\alpha_s - \alpha_t}{1 - \alpha_t}$$



# The model & The Loss

## Enter the Neural Network ( $\mu_\theta$ ):

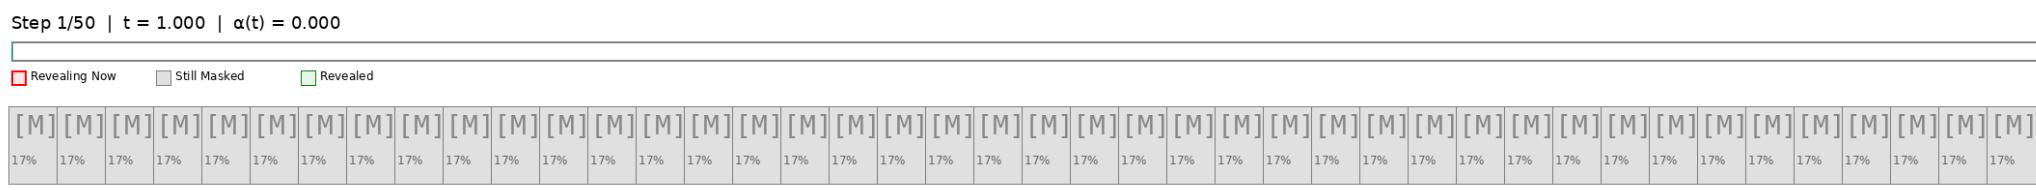
- During generation, we obviously don't know  $x_0$ .
  - Predicted categorical probabilities for the original token  $\mu_\theta(x_t, t) \approx P(x_0)$ .
  - **The Reverse Transition**  $q_\theta(x_s|x_t, \hat{x}_0)$ : We take the formula for the posterior (previous slide) and substitute the true  $x_0$  with our model's prediction  $\hat{x}_0$

## **Final Objective:**

$$\mathcal{L} = \int_0^1 w(t) \cdot \text{CE}(x_0, \mu_\theta(x_t, t)) dt$$

**Weighting:**  $w(t) = \frac{\alpha'(t)}{1-\alpha(t)}$ . The more tokens are masked, the less we weight the error.

**Training:** We sample a random  $t$ , mask the input accordingly, and compute the loss.



# Why Cross Entropy?

**The key idea:** Binary rule - *Stay Original or Become Mask.*

**ELBO:** minimize KL-Div between the true reverse step  $q(x_{t-1}|x_t, x_0)$  and the model's approximated step  $p_\theta(x_{t-1}|x_t)$

In masked diffusion, the "true" reverse step is effectively deterministic: if a token unmasks, it **must** reveal the specific ground truth token  $x_0$ .

The model is trained to predict the clean data  $\hat{x}_0$  directly, rather than predicting the slightly less noisy state  $x_{t-1}$ , aka “classification”.

**The Simplification:** The KL divergence between a deterministic distribution (the one-hot ground truth  $x_0$ ) and a predicted probability distribution is, by definition, Cross Entropy.

$$D_{\text{KL}}(\text{OneHot}(x_0) \parallel p_\theta(x)) = - \sum \mathbf{1}_{x=x_0} \log(p_\theta(x)) = \text{CrossEntropy}$$

# Summary

**Forward Process:**

$$q(x_t|x_s) = \frac{\alpha_t}{\alpha_s} \cdot \delta_{x_s} + (1 - \frac{\alpha_t}{\alpha_s}) \cdot \delta_{Mask}$$

**Reverse Process:**

$$q(x_s|x_t, x_0) \approx q_\theta(x_s|x_t, \hat{x}_0) = q_\theta(x_s|x_t, \mu_\theta(x_t, t))$$

$$P(\text{reveal}) = \frac{\alpha_s - \alpha_t}{1 - \alpha_t}$$

$\alpha_t$  is a masking schedule chosen by us

## From 1D to Multidimensional:

For e.g. text “multidimensional” refers to the sequence of tokens

**We assume the forward / reverse process is independent for each token.**

-> Simply vectorize the equations.

Note: The tokens are mathematically independent during noising, but the Neural Network mixes their information (via Self-Attention) during denoising.

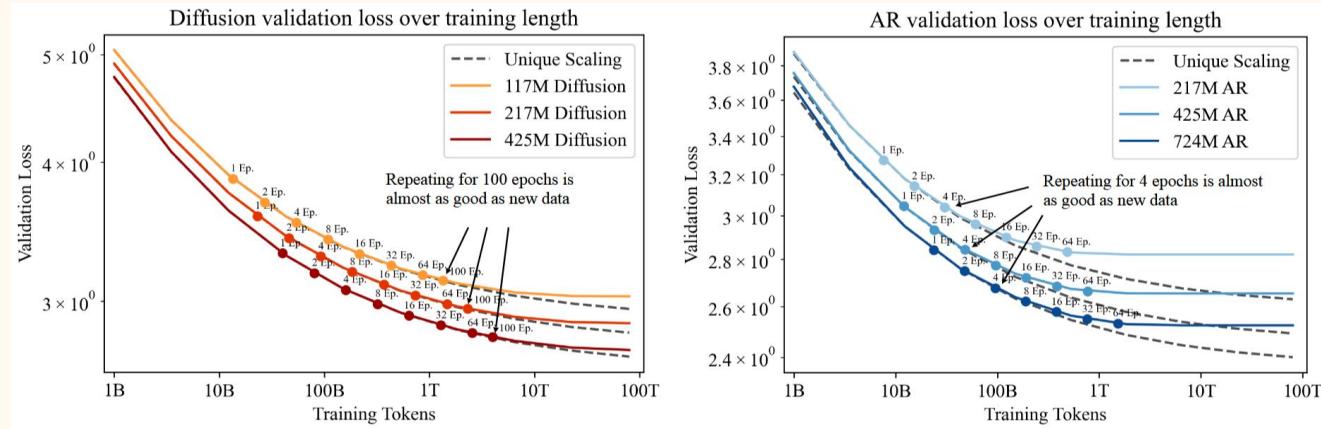
# State dependent masking schedules (GenMD4)

- Once a token is unmasked, it is frozen -> Commitment risk!
- Standard MD4 unmasks all tokens at the same speed

Aspect	Scalar Masking (MD4)	Vector Masking (GenMD4)
Core idea	One global noise schedule	One schedule per vocabulary item
$\alpha$ definition	$\alpha(t)$	$\alpha_i(t)$
Crucial formula	$\alpha(t) = (1 - 2\epsilon)f(t) + \varepsilon$	$\alpha_i(t) = 1 - (1 - \epsilon)t^{w_i}$
Learnable params	Model Params	Model Params + Vocab size
Schedule shape	Fixed (cosine / linear / poly)	Fixed family (poly), <b>learned speed</b>
Reduces to MD4	-	If all $w_i$ were equal

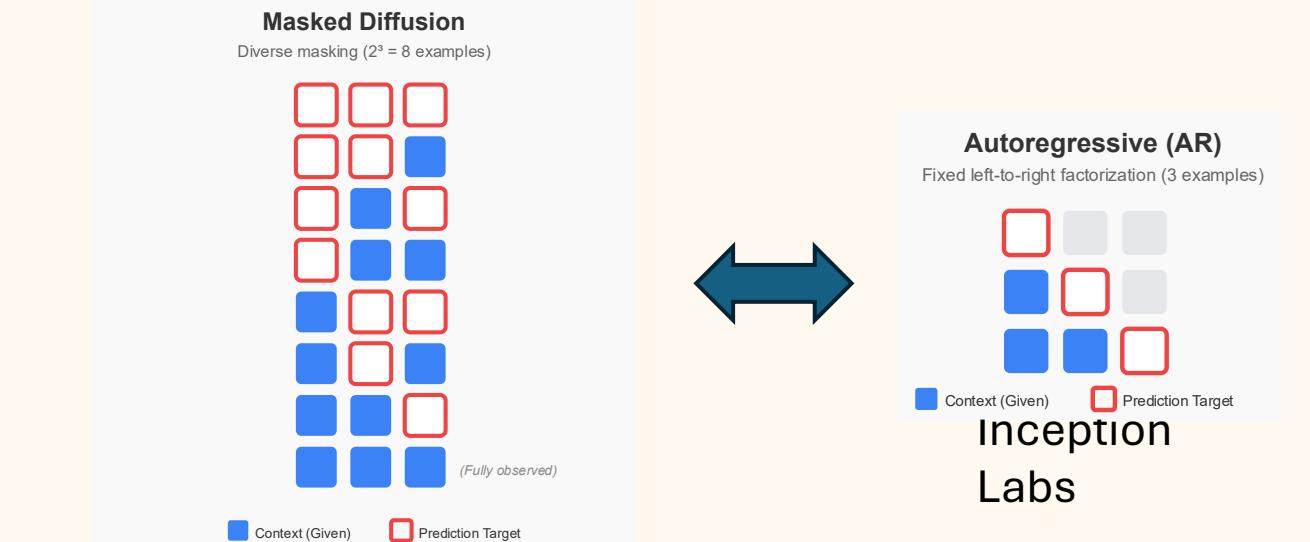
# Masked Diffusion vs Autoregressive Models

- Can be better or worse compared to AR
- Recent paper showed that:
  1. Masked diffusion models are more data efficient
  2. But, need longer training to achieve the same loss
- Data constrained? -> diffusion
- Compute constrained? -> AR



<https://arxiv.org/pdf/2507.15857v1.pdf>

- + Diffusion “sees and manipulates” across the whole sequence.
- + Thus, particularly good for tasks where this is useful (e.g., coding)
- + Faster for inference. You can pick the n steps.



# TDLM - „Tiny Discrete Language Model“

## My project: A tiny language model

- Original code in Jax ([github.com/google-deepmind/md4](https://github.com/google-deepmind/md4))
- Reimplemented in PyTorch from scratch

250 lines -&gt;

160 lines -&gt;

150 lines -&gt;

Name	Status	Änderungsdatum	Typ	Größe
model.py	🕒	12.12.2025 12:44	Python-Qu...	11 KB
train.ipynb	🕒	12.12.2025 15:17	Jupyter-Qu...	69 KB
transformer_impl.py	🕒	12.12.2025 12:44	Python-Qu...	7 KB
wine.txt	🕒	12.12.2025 12:50	Textdokum...	66.496 KB

## Model:

- 40M parameter transformer model
- Character level (vocab size is 90)

```
class MD4Config:
    """
    Specific parameter config used for this run.
    (Also includes the training params)
    """

    def __init__(self, vocab_size):
        # Model params
        self.vocab_size = vocab_size # (90)
        self.block_size = 128 # (sequence length)
        self.n_embd = 512
        self.n_head = 8
        self.n_layer = 8
        self.dropout = 0.1
        # Training params
        self.max_steps = 50000
        self.batch_size = 128
        self.learning_rate = 2e-4
        self.min_lr = 1e-5           # Decay to 5%
        self.warmup_steps = 1000
```

## Training:

- Dataset: Wine review, 60M tokens (<https://labelyourdata.com/datasets/wine-review-dataset>)
- Trained for 800M tokens on (12 epochs)
- Training finished in 1h

# TDLM - Training Progression

```
Step 0 | Loss: 573.8300 | LR: 2.00e-07 | Speed: 2936859 tok/s
-> Evaluating...
-> Train Loss: 588.0111 | Val Loss: 586.6883
```

```
--- Generating ---
JMzuc%
bV@TtKEzMg]vNdS@,-*K%.@RVF6j@Iff'TBBt79@H]OABiK@R24
w*UnAEv
Pj_)21J]NUx_`y:v5]CCE-loG#];QEN#:EacMB!DE4jT@Is-Ha)FVp&$,"M/
-----
```

```
Step 8000 | Loss: 167.8975 | LR: 1.91e-04 | Speed: 248548 tok/s
-> Evaluating...
-> Train Loss: 153.9978 | Val Loss: 151.6332
```

```
--- Generating ---
ch malablock struazed with acidity aromas of coctacize jam-aced tak not,s of a soft. Bright, good with burtery plump and cherry
-----
```

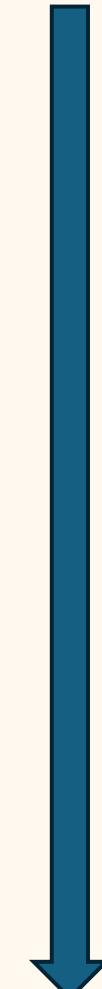
```
Step 8100 | Loss: 151.2366 | LR: 1.90e-04 | Speed: 158434 tok/s
```

```
Step 49000 | Loss: 123.0334 | LR: 1.02e-05 | Speed: 253801 tok/s
-> Evaluating...
-> Train Loss: 125.4636 | Val Loss: 124.3588
```

```
--- Generating ---
lemboy's much price Piorce the nose of cutes. Tobacco and smoke, spine line anilla red from a creamy note that's candy. This is
-----
```

```
Step 49100 | Loss: 142.4756 | LR: 1.02e-05 | Speed: 166285 tok/s
```

Loss / Steps



Now it knows fancy words 😊